


Model Knowledge Cutoffs vs. Research Paper Implementation Dates



FINDINGS:

- 5 out of 19 papers (26.3%) have at least one commit that predates some model's cutoff date.
- On average, 4.9% of model-paper pairs show paper first commits occurring before the model's knowledge cutoff date.

KEY FINDINGS & IMPLICATIONS

IMPLICATIONS:

- For 95.1% of model-paper pairs, models potentially had access to the research concepts during training.
- However, 4.9% of pairs show models potentially generating implementations for papers they were not trained on.
- This visualization helps identify which models/papers need contamination analysis.

This visualization tracks the temporal relationship between language model knowledge cutoff and research paper commit dates.

Percentages show: % of models with knowledge cutoff before each commit point